

Tidy Text Mining

Julia Silge and David Robinson
2016-07-06

Contents

Chapter 1

Introduction

This intro will be changed a lot to serve as a general and friendly intro to the topic.

1.1 What is tidy text?

As described by Hadley Wickham(?), tidy data has a specific structure:

- each variable is a column
- each observation is a row
- each type of observational unit is a table

Tidy data sets allow manipulation with a standard set of “tidy” tools, including popular packages such as `dplyr`(?), `ggplot2`(?), and `broom`(?). These tools do not yet, however, have the infrastructure to work fluently with text data and natural language processing tools. In developing this package, we provide functions and supporting data sets to allow conversion of text to and from tidy formats, and to switch seamlessly between tidy tools and existing text mining packages.

We define the tidy text format as being one-token-per-document-per-row, and provide functionality to tokenize by commonly used units of text including words, n-grams, and sentences. At the same time, the `tidytext` package doesn’t expect a user to keep text data in a tidy form at all times during an analysis. The package includes functions to `tidy` objects (see the `broom` package(?)) from popular text mining R packages such as `tm`(?) and `quanteda`(?). This allows, for example, a workflow with easy reading, filtering, and processing to be done using `dplyr` and other tidy tools, after which the data can be converted into a document-term matrix for machine learning applications. The models can then be re-converted into a tidy form for interpretation and visualization with `ggplot2`.

1.2 Outline

We start by introducing the tidy text format, and some of the ways `dplyr`, `tidyr` and `tidytext` allow informative analyses of this structure.

- **Chapter 2** outlines the tidy text format and the `unnest_tokens` function. It also introduces the `gutenbergr` and `janeaustenr` packages, which provide useful literary text datasets that we’ll use throughout this book.
- **Chapter 3** shows how to perform sentiment analysis on a tidy text dataset, using the `sentiments` dataset from `tidytext` and `inner_join` from `dplyr`
- **Chapter 4** describes the method of TF-IDF (term frequency times inverse document frequency), for identifying terms that are especially specific to a particular document. (Other document stuff in this chapter perhaps?)

Text won't be tidy at all stages of an analysis.

- **Chapter 5** introduces methods for tidying document-term matrices and Corpus objects from the `tm` and `quanteda` packages, as well as for casting tidy text datasets into those formats.
- **Chapter 6** introduces the concept of topic modeling, and uses the `tidy` method for interpreting and visualizing the output of the `topicmodels` package.
- **Chapter 7** (TODO) introduces tidying methods for the `glove` package, which offer an interface to word2vec models. (*These methods are still being implemented so this chapter is far from written!*)

We conclude with two tidy text analyses that bring together multiple text-mining approaches we've learned.

- **Chapter 8** demonstrates an application of a tidy text analysis on the Yelp restaurant review dataset. We show a few approaches to predicting a star rating from a review's text, and see how well sentiment analysis (from Chapter 3) does at this task.
- **Chapter 9** TODO: find at least one other in-depth exploration of text data. Optional but I think would conclude the book well.

Chapter 2

The Tidy Text Format

Intro text may go here about the one-token-per-document-per-row and about what is explored in the chapter.

2.1 The `unnest_tokens` function

```
text <- c("Because I could not stop for Death -",
         "He kindly stopped for me -",
         "The Carriage held but just Ourselves -",
         "and Immortality")
```

```
text
```

```
## [1] "Because I could not stop for Death -" "He kindly stopped for me -"
## [3] "The Carriage held but just Ourselves -" "and Immortality"
```

This is a typical character vector that we might want to analyze. In order to turn it into a tidy text dataset, we first need to put it into a data frame:

```
library(dplyr)
text_df <- data_frame(line = 1:4, text = text)
```

```
text_df
```

```
## # A tibble: 4 x 2
##   line      text
## * <int>    <chr>
## 1     1 Because I could not stop for Death -
## 2     2           He kindly stopped for me -
## 3     3 The Carriage held but just Ourselves -
## 4     4           and Immortality
```

Notice that this data frame isn't yet compatible with tidy tools. We can't filter out words or count which occur most frequently, since each row is made up of multiple combined tokens. We need to turn this into **one-token-per-document-per-row**.

To do this, we use tidytext's `unnest_tokens` function:

```
text_df %>%
  unnest_tokens(word, text)
```

```
## # A tibble: 20 x 2
```

```
##      line      word
##      <int>    <chr>
## 1         1 because
## 2         1      i
## 3         1   could
## 4         1    not
## 5         1   stop
## 6         1    for
## 7         1  death
## 8         2     he
## 9         2 kindly
## 10        2 stopped
## # ... with 10 more rows
```

We’ve now split each row so that there’s one token (word) in each row of the new data frame. Also notice:

- Other columns, such as the line number each word came from, are retained
- Punctuation has been stripped
- By default, `unnest_tokens` turns the tokens lowercase, which makes them easier to compare or combine with other datasets. (Use the `to_lower = FALSE` argument to turn off this behavior).

Having the text data in this format lets us manipulate, process, and visualize the text using the standard set of tidy tools; namely `dplyr`, `tidyr`, `ggplot2`, and `broom`.

2.2 Example: the works of Jane Austen

Let’s use the text of Jane Austen’s 6 completed, published novels from the `janeaustenr` package, and transform them into a tidy format. `janeaustenr` provides them as a one-row-per-line format:

```
library(janeaustenr)
library(dplyr)
library(stringr)

original_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenum = row_number(),
         chapter = cumsum(str_detect(text, regex("^chapter [\\divxlc]",
                                                ignore_case = TRUE)))) %>%
  ungroup()
```

```
original_books
```

```
## # A tibble: 73,422 x 4
##           text                book linenum chapter
##           <chr>              <fctr>   <int>   <int>
## 1 SENSE AND SENSIBILITY Sense & Sensibility         1         0
## 2                        Sense & Sensibility         2         0
## 3      by Jane Austen Sense & Sensibility         3         0
## 4                        Sense & Sensibility         4         0
## 5      (1811) Sense & Sensibility         5         0
## 6                        Sense & Sensibility         6         0
## 7                        Sense & Sensibility         7         0
## 8                        Sense & Sensibility         8         0
## 9                        Sense & Sensibility         9         0
## 10     CHAPTER 1 Sense & Sensibility        10         1
```



```
## # ... with 73,412 more rows
```

To work with this as a tidy dataset, we need to restructure it as **one-token-per-row** format. The `unnest_tokens` function is a way to convert a dataframe with a text column to be one-token-per-row:

```
library(tidytext)
tidy_books <- original_books %>%
  unnest_tokens(word, text)
```

```
tidy_books
```

```
## # A tibble: 725,054 x 4
##       book      linenumber chapter      word
##       <fctr>      <int>    <int>    <chr>
## 1 Sense & Sensibility      1      0 sense
## 2 Sense & Sensibility      1      0 and
## 3 Sense & Sensibility      1      0 sensibility
## 4 Sense & Sensibility      3      0 by
## 5 Sense & Sensibility      3      0 jane
## 6 Sense & Sensibility      3      0 austen
## 7 Sense & Sensibility      5      0 1811
## 8 Sense & Sensibility     10      1 chapter
## 9 Sense & Sensibility     10      1 1
## 10 Sense & Sensibility     13      1 the
## # ... with 725,044 more rows
```

This function uses the `tokenizers` package to separate each line into words. The default tokenizing is for words, but other options include characters, ngrams, sentences, lines, paragraphs, or separation around a regex pattern.

Now that the data is in one-word-per-row format, we can manipulate it with tidy tools like `dplyr`. We can remove stop words (kept in the tidytext dataset `stop_words`) with an `anti_join`.

```
data("stop_words")

tidy_books <- tidy_books %>%
  anti_join(stop_words)
```

We can also use `count` to find the most common words in all the books as a whole.

```
tidy_books %>%
  count(word, sort = TRUE)
```

```
## # A tibble: 13,914 x 2
##       word      n
##       <chr> <int>
## 1 miss  1855
## 2 time  1337
## 3 fanny   862
## 4 dear   822
## 5 lady   817
## 6 sir    806
## 7 day    797
## 8 emma   787
## 9 sister 727
## 10 house 699
## # ... with 13,904 more rows
```

For example, this allows us to visualize the popular words using `ggplot2`:

```
library(ggplot2)

tidy_books %>%
  count(word, sort = TRUE)

## # A tibble: 13,914 x 2
##   word      n
##   <chr> <int>
## 1 miss  1855
## 2 time  1337
## 3 fanny   862
## 4 dear   822
## 5 lady   817
## 6 sir    806
## 7 day    797
## 8 emma   787
## 9 sister  727
## 10 house  699
## # ... with 13,904 more rows
```

2.2.1 The gutenbergr package

TODO: Now that we've introduced the janeaustenr package, also include a brief intro to the gutenbergr package.

Chapter 3

Sentiment Analysis with Tidy Data

3.1 The sentiments dataset

The

```
sentiments
```

```
## # A tibble: 23,165 x 4
##       word sentiment lexicon score
##       <chr>      <chr>   <chr> <int>
## 1    abacus      trust    nrc   NA
## 2   abandon     fear    nrc   NA
## 3   abandon negative    nrc   NA
## 4   abandon sadness    nrc   NA
## 5  abandoned    anger    nrc   NA
## 6  abandoned    fear    nrc   NA
## 7  abandoned negative    nrc   NA
## 8  abandoned sadness    nrc   NA
## 9 abandonment  anger    nrc   NA
## 10 abandonment fear      nrc   NA
## # ... with 23,155 more rows
```

3.2 Sentiment analysis with inner join

Sentiment analysis can be done as an inner join. Three sentiment lexicons are in the tidytext package in the `sentiment` dataset. Let's look at the words with a joy score from the NRC lexicon. What are the most common joy words in *Emma*?

```
nrcjoy <- sentiments %>%
  filter(lexicon == "nrc", sentiment == "joy")

tidy_books %>%
  filter(book == "Emma") %>%
  semi_join(nrcjoy) %>%
  count(word, sort = TRUE)
```

```
## # A tibble: 298 x 2
##       word      n
##       <chr> <int>
```

```
## 1      friend  166
## 2       hope  143
## 3      happy  125
## 4       love  117
## 5       deal   92
## 6      found   92
## 7 happiness   76
## 8      pretty   68
## 9       true   66
## 10    comfort   65
## # ... with 288 more rows
```

Or instead we could examine how sentiment changes during each novel. Let's find a sentiment score for each word using the Bing lexicon, then count the number of positive and negative words in defined sections of each novel.


```
library(tidyr)
bing <- sentiments %>%
  filter(lexicon == "bing") %>%
  select(-score)

janeaustensentiment <- tidy_books %>%
  inner_join(bing) %>%
  count(book, index = linenumber %/% 80, sentiment) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative)
```

Now we can plot these sentiment scores across the plot trajectory of each novel.

```
library(ggplot2)

ggplot(janeaustensentiment, aes(index, sentiment, fill = book)) +
  geom_bar(stat = "identity", show.legend = FALSE) +
  facet_wrap(~book, ncol = 2, scales = "free_x")
```



3.2.1 Most common positive and negative words

One advantage of having the data frame with both sentiment and word is that we can analyze word counts that contribute to each sentiment.

```
bing_word_counts <- tidy_books %>%
  inner_join(bing) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()
```

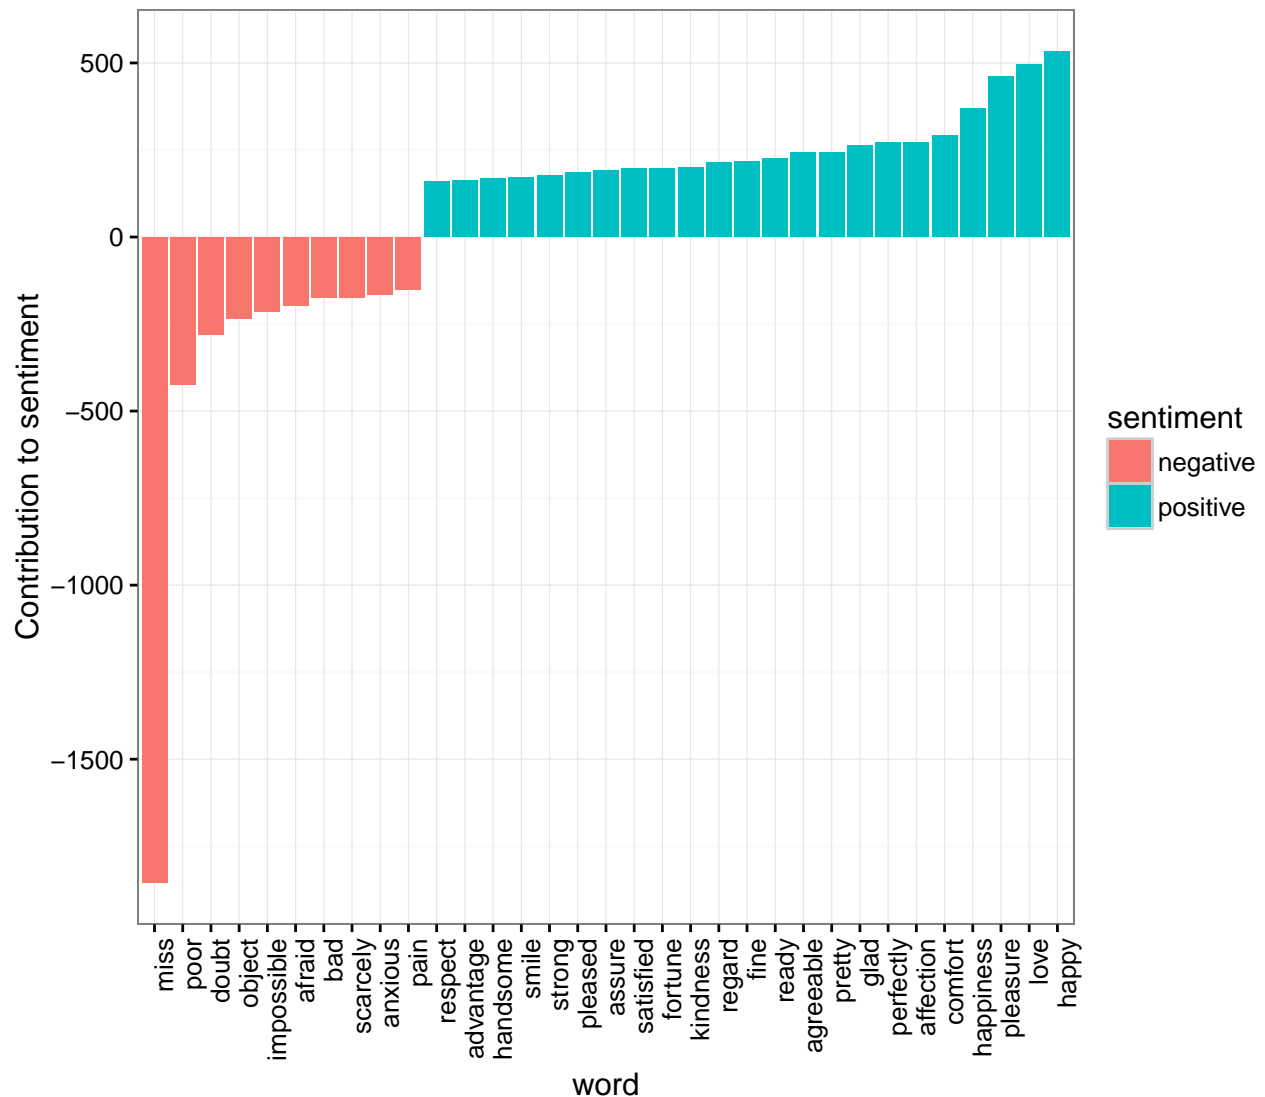
```
bing_word_counts
```

```
## # A tibble: 2,555 x 3
##       word sentiment      n
```

```
##      <chr>      <chr> <int>
## 1  abominable  negative   17
## 2  abominably  negative    7
## 3   abominate  negative    3
## 4    abound   positive    1
## 5    abrupt   negative    5
## 6  abruptly   negative   12
## 7   absence   negative  111
## 8    absurd   negative   19
## 9  absurdity   negative   12
## 10 abundance  positive   14
## # ... with 2,545 more rows
```

This can be shown visually, and we can pipe straight into `ggplot2` because of the way we are consistently using tools built for handling tidy data frames.

```
bing_word_counts %>%
  filter(n > 150) %>%
  mutate(n = ifelse(sentiment == "negative", -n, n)) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n, fill = sentiment)) +
  geom_bar(stat = "identity") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  ylab("Contribution to sentiment")
```



This lets us spot an anomaly in the sentiment analysis; the word “miss” is coded as negative but it is used as a title for young, unmarried women in Jane Austen’s works. If it were appropriate for our purposes, we could easily add “miss” to a custom stop-words list using `bind_rows`.

3.2.2 Wordclouds

We’ve seen that this tidy text mining approach works well with `ggplot2`, but having our data in a tidy format is useful for other plots as well.

For example, consider the `wordcloud` package. Let’s look at the most common words in Jane Austen’s works as a whole again.

```
library(wordcloud)

tidy_books %>%
  count(word) %>%
  with(wordcloud(word, n, max.words = 100))
```



```
austen_chapters <- austen_books() %>%
  group_by(book) %>%
  unnest_tokens(chapter, text, token = "regex", pattern = "Chapter|CHAPTER [\\dIVXLC]") %>%
  ungroup()
austen_chapters %>% group_by(book) %>% summarise(chapters = n())
```

```
## # A tibble: 6 x 2
##           book chapters
##           <fctr>   <int>
## 1 Sense & Sensibility    51
## 2 Pride & Prejudice      62
## 3 Mansfield Park        49
## 4 Emma                  56
## 5 Northanger Abbey      32
## 6 Persuasion            25
```

We have recovered the correct number of chapters in each novel (plus an “extra” row for each novel title). In this data frame, each row corresponds to one chapter.

Near the beginning of this vignette, we used a similar regex to find where all the chapters were in Austen’s novels for a tidy data frame organized by one-word-per-row. We can use tidy text analysis to ask questions such as what are the most negative chapters in each of Jane Austen’s novels? First, let’s get the list of negative words from the Bing lexicon. Second, let’s make a dataframe of how many words are in each chapter so we can normalize for the length of chapters. Then, let’s find the number of negative words in each chapter and divide by the total words in each chapter. Which chapter has the highest proportion of negative words?

```
bingnegative <- sentiments %>%
  filter(lexicon == "bing", sentiment == "negative")

wordcounts <- tidy_books %>%
  group_by(book, chapter) %>%
  summarize(words = n())

tidy_books %>%
  semi_join(bingnegative) %>%
  group_by(book, chapter) %>%
  summarize(negativewords = n()) %>%
  left_join(wordcounts, by = c("book", "chapter")) %>%
  mutate(ratio = negativewords/words) %>%
  filter(chapter != 0) %>%
  top_n(1)
```

```
## Source: local data frame [6 x 5]
## Groups: book [6]
##
##           book chapter negativewords words      ratio
##           (fctr)   (int)         (int) (int)      (dbl)
## 1 Sense & Sensibility    29           172  1135 0.1515419
## 2 Pride & Prejudice      34           108   646 0.1671827
## 3 Mansfield Park        45           132   884 0.1493213
## 4 Emma                  15           147  1012 0.1452569
## 5 Northanger Abbey      27            55   337 0.1632047
## 6 Persuasion            21           215  1948 0.1103696
```

These are the chapters with the most negative words in each book, normalized for number of words in the

chapter. What is happening in these chapters? In Chapter 29 of *Sense and Sensibility* Marianne finds out what an awful person Willoughby is by letter, and in Chapter 34 of *Pride and Prejudice* Mr. Darcy proposes for the first time (so badly!). Chapter 45 of *Mansfield Park* is almost the end, when Tom is sick with consumption and Mary is revealed as mercenary and uncaring, Chapter 15 of *Emma* is when horrifying Mr. Elton proposes, and Chapter 27 of *Northanger Abbey* is a short chapter where Catherine gets a terrible letter from her inconstant friend Isabella. Chapter 21 of *Persuasion* is when Anne's friend tells her all about Mr. Elliott's immoral past.

Chapter 4

TF-IDF: Analyzing word and document frequency

A central question in text mining and natural language processing is how to quantify what a document is about. Can we do this by looking at the words that make up the document? One measure of how important a word may be is its *term frequency* (tf), how frequently a word occurs in a document. There are words in a document, however, that occur many times but may not be important; in English, these are probably words like “the”, “is”, “of”, and so forth. We might take the approach of adding words like these to a list of stop words and removing them before analysis, but it is possible that some of these words might be more important in some documents than others. A list of stop words is not a sophisticated approach to adjusting term frequency for commonly used words.

4.1 Term frequency and inverse document frequency

Another approach is to look at a term’s *inverse document frequency* (idf), which decreases the weight for commonly used words and increases the weight for words that are not used very much in a collection of documents. This can be combined with term frequency to calculate a term’s *tf-idf*, the frequency of a term adjusted for how rarely it is used. It is intended to measure how important a word is to a document in a collection (or corpus) of documents. It is a rule-of-thumb or heuristic quantity; while it has proved useful in text mining, search engines, etc., its theoretical foundations are considered less than firm by information theory experts. The inverse document frequency for any given term is defined as

$$idf(\text{term}) = \ln \left(\frac{n_{\text{documents}}}{n_{\text{documents containing term}}} \right)$$

We can use tidy data principles, as described in the main vignette, to approach tf-idf analysis and use consistent, effective tools to quantify how important various terms are in a document that is part of a collection.

Let’s look at the published novels of Jane Austen and examine first term frequency, then tf-idf. We can start just by using dplyr verbs such as `group_by` and `join`. What are the most commonly used words in Jane Austen’s novels? (Let’s also calculate the total words in each novel here, for later use.)

```
library(dplyr)
library(janeaustenr)
library(tidytext)
book_words <- austen_books() %>%
  unnest_tokens(word, text) %>%
```

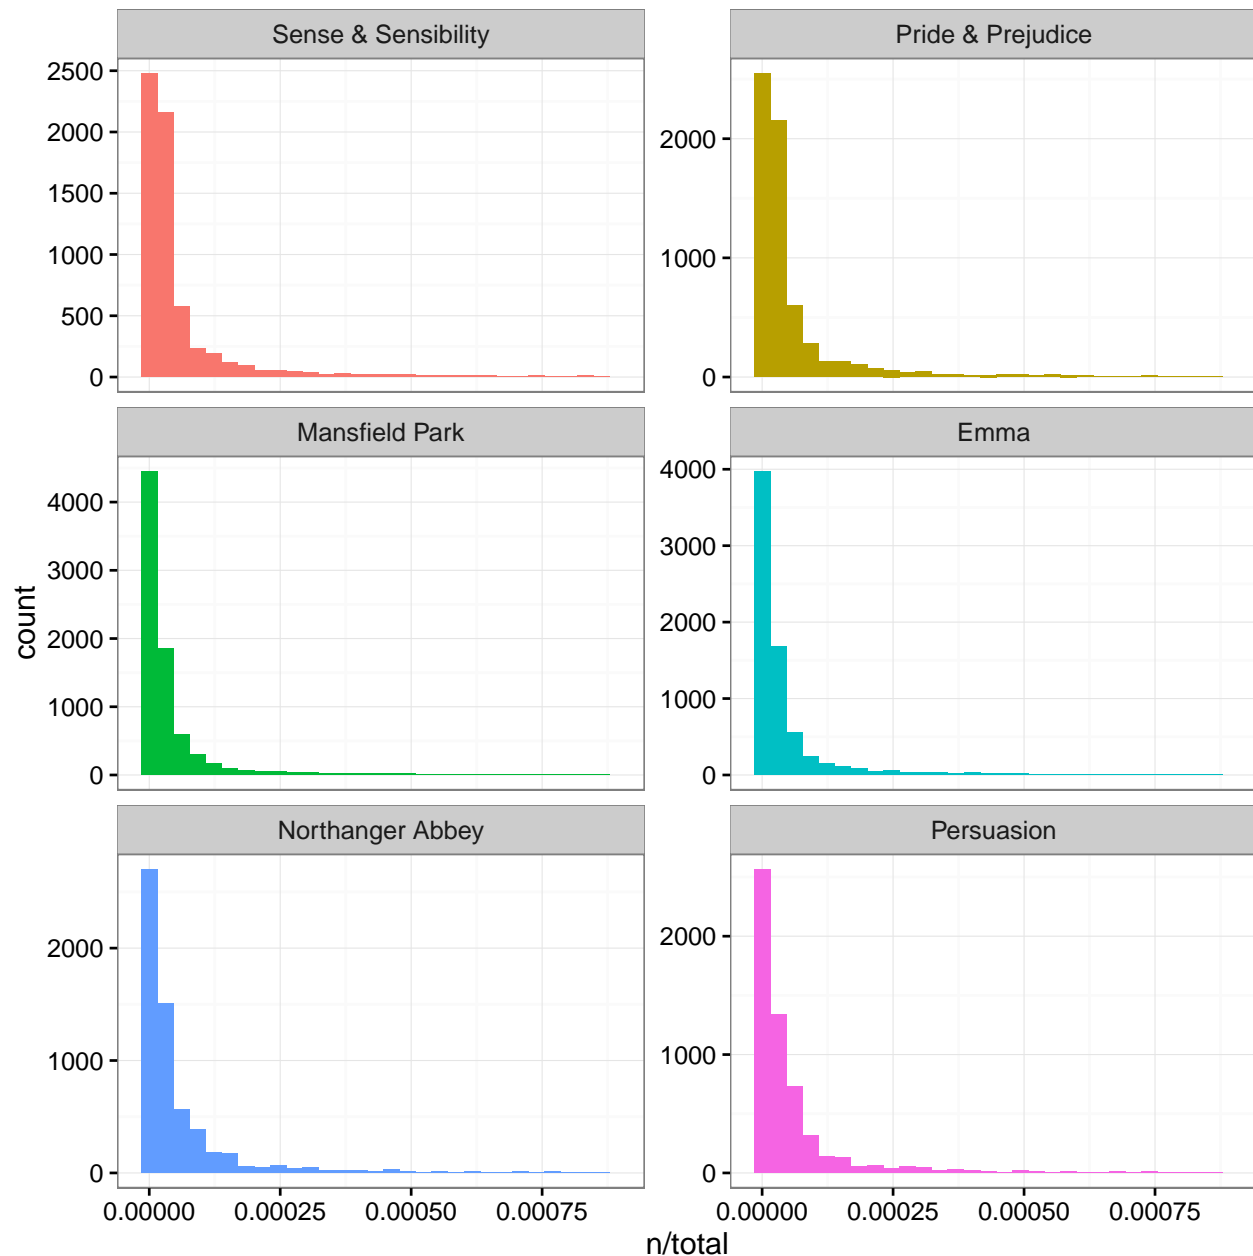
```
count(book, word, sort = TRUE) %>%
ungroup()

total_words <- book_words %>% group_by(book) %>% summarize(total = sum(n))
book_words <- left_join(book_words, total_words)
book_words
```

```
## # A tibble: 40,379 x 4
##           book word      n total
##           <fctr> <chr> <int> <int>
## 1 Sense & Sensibility to 4116 119957
## 2 Sense & Sensibility the 4105 119957
## 3 Sense & Sensibility of 3571 119957
## 4 Sense & Sensibility and 3490 119957
## 5 Sense & Sensibility her 2543 119957
## 6 Sense & Sensibility a 2092 119957
## 7 Sense & Sensibility i 1998 119957
## 8 Sense & Sensibility in 1979 119957
## 9 Sense & Sensibility was 1861 119957
## 10 Sense & Sensibility it 1755 119957
## # ... with 40,369 more rows
```

The usual suspects are here, “the”, “and”, “to”, and so forth. Let’s look at the distribution of n/total for each novel, the number of times a word appears in a novel divided by the total number of terms (words) in that novel. This is exactly what term frequency is.

```
library(ggplot2)
ggplot(book_words, aes(n/total, fill = book)) +
  geom_histogram(show.legend = FALSE) +
  xlim(NA, 0.0009) +
  facet_wrap(~book, ncol = 2, scales = "free_y")
```



There are very long tails to the right for these novels (those extremely common words!) that we have not shown in these plots. These plots exhibit similar distributions for all the novels, with many words that occur rarely and fewer words that occur frequently.

4.2 The bind_tf_idf function

The idea of tf-idf is to find the important words for the content of each document by decreasing the weight for commonly used words and increasing the weight for words that are not used very much in a collection or corpus of documents, in this case, the group of Jane Austen's novels as a whole. Calculating tf-idf attempts to find the words that are important (i.e., common) in a text, but not *too* common. Let's do that now.

```
book_words <- book_words %>%
  bind_tf_idf(word, book, n)
```

```
book_words
```

```
## # A tibble: 40,379 x 7
##       book word      n total      tf      idf tf_idf
##       <fctr> <chr> <int> <int>    <dbl> <dbl> <dbl>
## 1 Sense & Sensibility to 4116 119957 0.03431230 0 0
## 2 Sense & Sensibility the 4105 119957 0.03422060 0 0
## 3 Sense & Sensibility of 3571 119957 0.02976900 0 0
## 4 Sense & Sensibility and 3490 119957 0.02909376 0 0
## 5 Sense & Sensibility her 2543 119957 0.02119926 0 0
## 6 Sense & Sensibility a 2092 119957 0.01743958 0 0
## 7 Sense & Sensibility i 1998 119957 0.01665597 0 0
## 8 Sense & Sensibility in 1979 119957 0.01649758 0 0
## 9 Sense & Sensibility was 1861 119957 0.01551389 0 0
## 10 Sense & Sensibility it 1755 119957 0.01463024 0 0
## # ... with 40,369 more rows
```

Notice that idf and thus tf-idf are zero for these extremely common words. These are all words that appear in all six of Jane Austen's novels, so the idf term (which will then be the natural log of 1) is zero. The inverse document frequency (and thus tf-idf) is very low (near zero) for words that occur in many of the documents in a collection; this is how this approach decreases the weight for common words. The inverse document frequency will be a higher number for words that occur in fewer of the documents in the collection. Let's look at terms with high tf-idf in Jane Austen's works.

```
book_words %>%
```

```
  select(-total) %>%
  arrange(desc(tf_idf))
```

```
## # A tibble: 40,379 x 6
##       book word      n      tf      idf      tf_idf
##       <fctr> <chr> <int>    <dbl>    <dbl>    <dbl>
## 1 Sense & Sensibility elinor 623 0.005193528 1.791759 0.009305552
## 2 Sense & Sensibility marianne 492 0.004101470 1.791759 0.007348847
## 3 Mansfield Park crawford 493 0.003072417 1.791759 0.005505032
## 4 Pride & Prejudice darcy 373 0.003052273 1.791759 0.005468939
## 5 Persuasion elliot 254 0.003036207 1.791759 0.005440153
## 6 Emma emma 786 0.004882109 1.098612 0.005363545
## 7 Northanger Abbey tilney 196 0.002519928 1.791759 0.004515105
## 8 Emma weston 389 0.002416209 1.791759 0.004329266
## 9 Pride & Prejudice bennet 294 0.002405813 1.791759 0.004310639
## 10 Persuasion wentworth 191 0.002283132 1.791759 0.004090824
## # ... with 40,369 more rows
```

Here we see all proper nouns, names that are in fact important in these novels. None of them occur in all of novels, and they are important, characteristic words for each text. Some of the values for idf are the same for different terms because there are 6 documents in this corpus and we are seeing the numerical value for $\ln(6/1)$, $\ln(6/2)$, etc. Let's look specifically at *Pride and Prejudice*.

```
book_words %>%
```

```
  filter(book == "Pride & Prejudice") %>%
  select(-total) %>%
  arrange(desc(tf_idf))
```

```
## # A tibble: 6,538 x 6
##       book word      n      tf      idf      tf_idf
##       <fctr> <chr> <int>    <dbl>    <dbl>    <dbl>
## 1 Pride & Prejudice darcy 373 0.0030522732 1.7917595 0.005468939
```



```
## 2 Pride & Prejudice bennet 294 0.0024058132 1.7917595 0.004310639
## 3 Pride & Prejudice bingley 257 0.0021030408 1.7917595 0.003768143
## 4 Pride & Prejudice elizabeth 597 0.0048852738 0.6931472 0.003386214
## 5 Pride & Prejudice wickham 162 0.0013256522 1.7917595 0.002375250
## 6 Pride & Prejudice collins 156 0.0012765540 1.7917595 0.002287278
## 7 Pride & Prejudice lydia 133 0.0010883441 1.7917595 0.001950051
## 8 Pride & Prejudice lizzy 95 0.0007773886 1.7917595 0.001392893
## 9 Pride & Prejudice longbourn 88 0.0007201074 1.7917595 0.001290259
## 10 Pride & Prejudice gardiner 84 0.0006873752 1.7917595 0.001231611
## # ... with 6,528 more rows
```

These words are, as measured by tf-idf, the most important to *Pride and Prejudice* and most readers would likely agree.

Chapter 5

Tidying and casting document-term matrices

Intro text here.

5.1 Tidying a document-term matrix

Many existing text mining datasets are in the form of a `DocumentTermMatrix` class (from the `tm` package). For example, consider the corpus of 2246 Associated Press articles from the `topicmodels` package:

```
library(tm)
data("AssociatedPress", package = "topicmodels")
AssociatedPress

## <<DocumentTermMatrix (documents: 2246, terms: 10473)>>
## Non-/sparse entries: 302031/23220327
## Sparsity          : 99%
## Maximal term length: 18
## Weighting         : term frequency (tf)
```

If we want to analyze this with tidy tools, we need to turn it into a one-token-per-document-per-row data frame first. The `tidy` function does this. (For more on the `tidy` verb, see the `broom` package).

```
library(dplyr)
library(tidytext)

ap_td <- tidy(AssociatedPress)
```

Just as shown in this vignette, having the text in this format is convenient for analysis with the `tidytext` package. For example, you can perform sentiment analysis on these newspaper articles.

```
bing <- sentiments %>%
  filter(lexicon == "bing") %>%
  select(word, sentiment)

ap_sentiments <- ap_td %>%
  inner_join(bing, by = c(term = "word"))

ap_sentiments
```

```
## # A tibble: 30,094 x 4
##   document    term count sentiment
##   <int>    <chr> <dbl>    <chr>
## 1         1 assault     1  negative
## 2         1 complex     1  negative
## 3         1  death     1  negative
## 4         1   died     1  negative
## 5         1   good     2  positive
## 6         1 illness     1  negative
## 7         1 killed     2  negative
## 8         1   like     2  positive
## 9         1  liked     1  positive
## 10        1 miracle     1  positive
## # ... with 30,084 more rows
```

We can find the most negative documents:

```
library(tidyr)

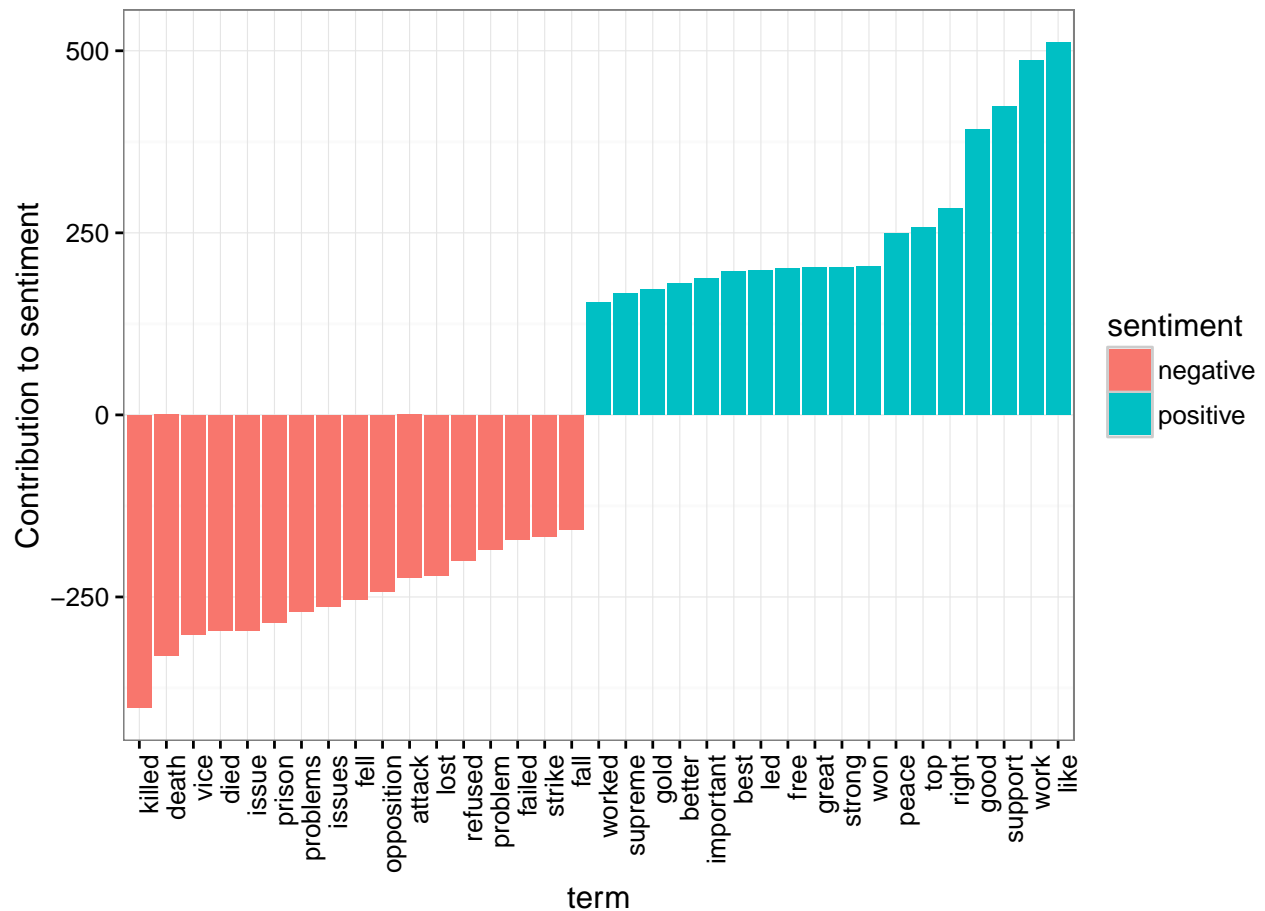
ap_sentiments %>%
  count(document, sentiment, wt = count) %>%
  ungroup() %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative) %>%
  arrange(sentiment)
```

```
## # A tibble: 2,190 x 4
##   document negative positive sentiment
##   <int>    <dbl>    <dbl>    <dbl>
## 1    1251      54        6      -48
## 2    1380      53        5      -48
## 3     531      51        9      -42
## 4      43      45       11      -34
## 5    1263      44       10      -34
## 6    2178      40        6      -34
## 7     334      45       12      -33
## 8    1664      38        5      -33
## 9    2147      47       14      -33
## 10     516      38        6      -32
## # ... with 2,180 more rows
```

Or visualize which words contributed to positive and negative sentiment:

```
library(ggplot2)

ap_sentiments %>%
  count(sentiment, term, wt = count) %>%
  ungroup() %>%
  filter(n >= 150) %>%
  mutate(n = ifelse(sentiment == "negative", -n, n)) %>%
  mutate(term = reorder(term, n)) %>%
  ggplot(aes(term, n, fill = sentiment)) +
  geom_bar(stat = "identity") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  ylab("Contribution to sentiment")
```



Note that a tidier is also available for the `dfm` class from the `quanteda` package:

```
data("inaugCorpus", package = "quanteda")
d <- quanteda::dfm(inaugCorpus)
```

```
## Creating a dfm from a corpus ...
##   ... lowercasing
##   ... tokenizing
##   ... indexing documents: 57 documents
##   ... indexing features: 9,214 feature types
##   ... created a 57 x 9215 sparse dfm
##   ... complete.
## Elapsed time: 0.131 seconds.
```

```
d
```

```
## Document-feature matrix of: 57 documents, 9,215 features.
```

```
tidy(d)
```

```
## # A tibble: 43,719 x 3
##   document      term count
## *   <chr>      <chr> <dbl>
## 1 1789-Washington fellow-citizens 1
## 2 1797-Adams fellow-citizens 3
## 3 1801-Jefferson fellow-citizens 2
## 4 1809-Madison fellow-citizens 1
## 5 1813-Madison fellow-citizens 1
```

```
## 6      1817-Monroe fellow-citizens      5
## 7      1821-Monroe fellow-citizens      1
## 8      1841-Harrison fellow-citizens    11
## 9      1845-Polk fellow-citizens        1
## 10     1849-Taylor fellow-citizens       1
## # ... with 43,709 more rows
```

5.2 Casting tidy text data into a DocumentTermMatrix

Some existing text mining tools or algorithms work only on sparse document-term matrices. Therefore, `tidytext` provides `cast_` verbs for converting from a tidy form to these matrices.

```
ap_td
```

```
## # A tibble: 302,031 x 3
##   document      term count
## *   <int>      <chr> <dbl>
## 1         1    adding     1
## 2         1    adult     2
## 3         1      ago     1
## 4         1  alcohol     1
## 5         1 allegedly     1
## 6         1    allen     1
## 7         1 apparently     2
## 8         1   appeared     1
## 9         1   arrested     1
## 10        1   assault     1
## # ... with 302,021 more rows
```

```
# cast into a Document-Term Matrix
```

```
ap_td %>%
  cast_dtm(document, term, count)
```

```
## <<DocumentTermMatrix (documents: 2246, terms: 10473)>>
## Non-/sparse entries: 302031/23220327
## Sparsity           : 99%
## Maximal term length: 18
## Weighting           : term frequency (tf)
```

```
# cast into a Term-Document Matrix
```

```
ap_td %>%
  cast_tdm(term, document, count)
```

```
## <<TermDocumentMatrix (terms: 10473, documents: 2246)>>
##Non-/sparse entries: 302031/23220327
## Sparsity           : 99%
## Maximal term length: 18
## Weighting           : term frequency (tf)
```

```
# cast into quanteda's dfm
```

```
ap_td %>%
  cast_dfm(term, document, count)
```

```
## Document-feature matrix of: 10,473 documents, 2,246 features.
```

```
# cast into a Matrix object
```

```
m <- ap_td %>%
```

```
cast_sparse(document, term, count)
class(m)
```

```
## [1] "dgCMatrix"
## attr(,"package")
## [1] "Matrix"
```

```
dim(m)
```

```
## [1] 2246 10473
```

This allows for easy reading, filtering, and processing to be done using dplyr and other tidy tools, after which the data can be converted into a document-term matrix for machine learning applications.

5.3 Tidying corpus objects with metadata

You can also tidy Corpus objects from the tm package. For example, consider a Corpus containing 20 documents:

```
reut21578 <- system.file("texts", "crude", package = "tm")
reuters <- VCorpus(DirSource(reut21578),
  readerControl = list(reader = readReut21578XMLasPlain))
```

```
reuters
```

```
## <<VCorpus>>
## Metadata: corpus specific: 0, document level (indexed): 0
## Content: documents: 20
```

The tidy verb creates a table with one row per document:

```
reuters_td <- tidy(reuters)
reuters_td
```

```
## # A tibble: 20 x 17
##               author      datetimestamp description
##               <chr>          <time>         <chr>
## 1               <NA> 1987-02-26 12:00:56
## 2 BY TED D'AFFLISIO, Reuters 1987-02-26 12:34:11
## 3               <NA> 1987-02-26 13:18:00
## 4               <NA> 1987-02-26 13:21:01
## 5               <NA> 1987-02-26 14:00:57
## 6               <NA> 1987-02-28 22:25:46
## 7 By Jeremy Clift, Reuters 1987-02-28 22:39:14
## 8               <NA> 1987-03-01 00:27:27
## 9               <NA> 1987-03-01 03:22:30
## 10              <NA> 1987-03-01 13:31:44
## 11              <NA> 1987-03-01 20:05:49
## 12              <NA> 1987-03-02 02:39:23
## 13              <NA> 1987-03-02 02:43:22
## 14              <NA> 1987-03-02 02:43:41
## 15              <NA> 1987-03-02 03:25:42
## 16              <NA> 1987-03-02 06:20:05
## 17              <NA> 1987-03-02 06:28:26
## 18              <NA> 1987-03-02 07:13:46
## 19 By BERNICE NAPACH, Reuters 1987-03-02 09:38:34
```

```
## 20          <NA> 1987-03-02 09:49:06
##          heading      id language      origin
##          <chr> <chr>    <chr>      <chr>
## 1      DIAMOND SHAMROCK (DIA) CUTS CRUDE PRICES 127      en Reuters-21578 XML
## 2      OPEC MAY HAVE TO MEET TO FIRM PRICES - ANALYSTS 144      en Reuters-21578 XML
## 3      TEXACO CANADA <TXC> LOWERS CRUDE POSTINGS 191      en Reuters-21578 XML
## 4      MARATHON PETROLEUM REDUCES CRUDE POSTINGS 194      en Reuters-21578 XML
## 5      HOUSTON OIL <HO> RESERVES STUDY COMPLETED 211      en Reuters-21578 XML
## 6      KUWAIT SAYS NO PLANS FOR EMERGENCY OPEC TALKS 236      en Reuters-21578 XML
## 7      INDONESIA SEEN AT CROSSROADS OVER ECONOMIC CHANGE 237      en Reuters-21578 XML
## 8      SAUDI RIYAL DEPOSIT RATES REMAIN FIRM 242      en Reuters-21578 XML
## 9      QATAR UNVEILS BUDGET FOR FISCAL 1987/88 246      en Reuters-21578 XML
## 10     SAUDI ARABIA REITERATES COMMITMENT TO OPEC PACT 248      en Reuters-21578 XML
## 11     SAUDI FEBRUARY CRUDE OUTPUT PUT AT 3.5 MLN BPD 273      en Reuters-21578 XML
## 12     GULF ARAB DEPUTY OIL MINISTERS TO MEET IN BAHRAIN 349      en Reuters-21578 XML
## 13     SAUDI ARABIA REITERATES COMMITMENT TO OPEC ACCORD 352      en Reuters-21578 XML
## 14     KUWAIT MINISTER SAYS NO EMERGENCY OPEC TALKS SET 353      en Reuters-21578 XML
## 15     PHILADELPHIA PORT CLOSED BY TANKER CRASH 368      en Reuters-21578 XML
## 16     STUDY GROUP URGES INCREASED U.S. OIL RESERVES 489      en Reuters-21578 XML
## 17     STUDY GROUP URGES INCREASED U.S. OIL RESERVES 502      en Reuters-21578 XML
## 18     UNOCAL <UCL> UNIT CUTS CRUDE OIL POSTED PRICES 543      en Reuters-21578 XML
## 19     NYMEX WILL EXPAND OFF-HOUR TRADING APRIL ONE 704      en Reuters-21578 XML
## 20     ARGENTINE OIL PRODUCTION DOWN IN JANUARY 1987 708      en Reuters-21578 XML
## # ... with 10 more variables: topics <chr>, lewissplit <chr>, cgisplit <chr>, oldid <chr>,
## #   topics_cat <list>, places <list>, people <chr>, orgs <chr>, exchanges <chr>, text <chr>
```

Similarly, you can tidy a corpus object from the quanteda package:

```
library(quanteda)

data("inaugCorpus")

inaugCorpus
```

```
## Corpus consisting of 57 documents and 3 docvars.
```

```
inaug_td <- tidy(inaugCorpus)
inaug_td
```

```
## # A tibble: 57 x 4
```

```
##
```

```
##
```

```
## 1 Fellow-Citizens of the Senate and of the House of Representatives:\n\nAmong the vicissitudes incident t
## 2 Fellow citizens, I am again called upon by the voice of my country to execute the functions of its Chie
## 3 When it was first perceived, in early times, that no middle course for America remained between unlimi
## 4 Friends and Fellow Citizens:\n\nCalled upon to undertake the duties of the first executive office of ou
## 5 Proceeding, fellow citizens, to that qualification which the Constitution requires before my entrance
## 6 Unwilling to depart from examples of the most revered authority, I avail myself of the occasion now pre
## 7 About to add the solemnity of an oath to the obligations imposed by a second call to the station in whic
## 8 I should be destitute of feeling if I was not deeply affected by the strong proof which my fellow-citiz
## 9 Fellow citizens, I shall not attempt to describe the grateful emotions which the new and very distingui
## 10 In compliance with an usage coeval with the existence of our Federal Constitution, and sanctioned by t
## # ... with 47 more rows, and 3 more variables: Year <int>, President <chr>, FirstName <chr>
```

This lets us work with tidy tools like `unnest_tokens` to analyze the text alongside the metadata.


```

inaug_words <- inaug_td %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words)

inaug_words

## # A tibble: 49,621 x 4
##   Year President FirstName      word
##   <int>   <chr>   <chr>    <chr>
## 1  2013    Obama   Barack   waves
## 2  2013    Obama   Barack realizes
## 3  2013    Obama   Barack philadelphia
## 4  2013    Obama   Barack    400
## 5  2013    Obama   Barack    40
## 6  2013    Obama   Barack absolutism
## 7  2013    Obama   Barack  contour
## 8  2013    Obama   Barack newtown
## 9  2013    Obama   Barack   lanes
## 10 2013    Obama   Barack appalachia
## # ... with 49,611 more rows

```

We could then, for example, see how the appearance of a word changes over time:

```

inaug_freq <- inaug_words %>%
  count(Year, word) %>%
  ungroup() %>%
  complete(Year, word, fill = list(n = 0)) %>%
  group_by(Year) %>%
  mutate(year_total = sum(n),
         percent = n / year_total) %>%
  ungroup()

inaug_freq

```

```

## # A tibble: 490,200 x 5
##   Year      word      n year_total  percent
##   <int>   <chr> <dbl>    <dbl>    <dbl>
## 1  1789         1      0      529 0.000000000
## 2  1789    1,000      0      529 0.000000000
## 3  1789     100      0      529 0.000000000
## 4  1789 100,000,000      0      529 0.000000000
## 5  1789 120,000,000      0      529 0.000000000
## 6  1789     125      0      529 0.000000000
## 7  1789      13      0      529 0.000000000
## 8  1789    14th      1      529 0.001890359
## 9  1789    15th      0      529 0.000000000
## 10 1789     16      0      529 0.000000000
## # ... with 490,190 more rows

```

For example, we can use the broom package to perform logistic regression on each word.

```

models <- inaug_freq %>%
  group_by(word) %>%
  filter(sum(n) > 50) %>%
  do(tidy(glm(cbind(n, year_total - n) ~ Year, .,
              family = "binomial"))) %>%

```

```

ungroup() %>%
  filter(term == "Year")

models

## # A tibble: 113 x 6
##       word term      estimate std.error statistic    p.value
##       <chr> <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1      act Year  0.006894234 0.002191596  3.1457591 1.656564e-03
## 2    action Year  0.001634417 0.001959204  0.8342250 4.041542e-01
## 3 administration Year -0.006979577 0.001882474 -3.7076616 2.091819e-04
## 4    america Year  0.018890081 0.001584306 11.9232506 8.954525e-33
## 5   american Year  0.007084142 0.001321897  5.3590709 8.365105e-08
## 6  americans Year  0.032657656 0.003659114  8.9250184 4.456252e-19
## 7  authority Year -0.005640373 0.002336159 -2.4143787 1.576207e-02
## 8   business Year  0.003745929 0.002016455  1.8576801 6.321445e-02
## 9    called Year -0.001935068 0.002088388 -0.9265844 3.541423e-01
## 10  century Year  0.016480566 0.002495844  6.6032027 4.023687e-11
## # ... with 103 more rows

models %>%
  filter(term == "Year") %>%
  arrange(desc(abs(estimate)))

```

```

## # A tibble: 113 x 6
##       word term      estimate std.error statistic    p.value
##       <chr> <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1  americans Year  0.03265766 0.003659114  8.925018 4.456252e-19
## 2    america Year  0.01889008 0.001584306 11.923251 8.954525e-33
## 3  century Year  0.01648057 0.002495844  6.603203 4.023687e-11
## 4     live Year  0.01448914 0.002490610  5.817506 5.973212e-09
## 5 democracy Year  0.01432438 0.002394738  5.981606 2.209489e-09
## 6      god Year  0.01402582 0.001921362  7.299935 2.879058e-13
## 7 freedom Year  0.01366336 0.001320242 10.349129 4.223092e-25
## 8  foreign Year -0.01364998 0.002058045 -6.632497 3.300543e-11
## 9   earth Year  0.01303351 0.002291996  5.686532 1.296449e-08
## 10   world Year  0.01233715 0.001000739 12.328042 6.398240e-35
## # ... with 103 more rows

```

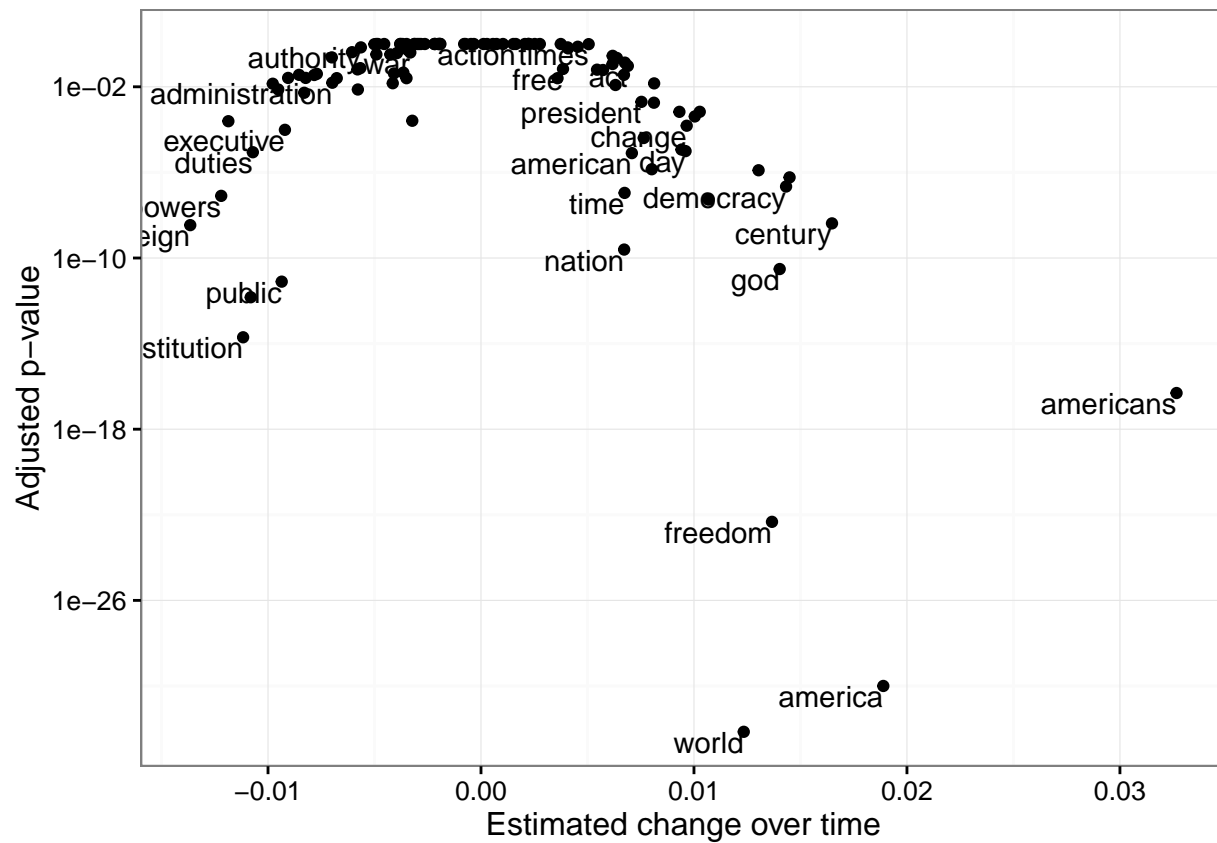
You can show these models as a volcano plot, which compares the effect size with the significance:

```

library(ggplot2)
theme_set(theme_bw())

models %>%
  mutate(adjusted.p.value = p.adjust(p.value)) %>%
  ggplot(aes(estimate, adjusted.p.value)) +
  geom_point() +
  scale_y_log10() +
  geom_text(aes(label = word), vjust = 1, hjust = 1,
            check_overlap = TRUE) +
  xlab("Estimated change over time") +
  ylab("Adjusted p-value")

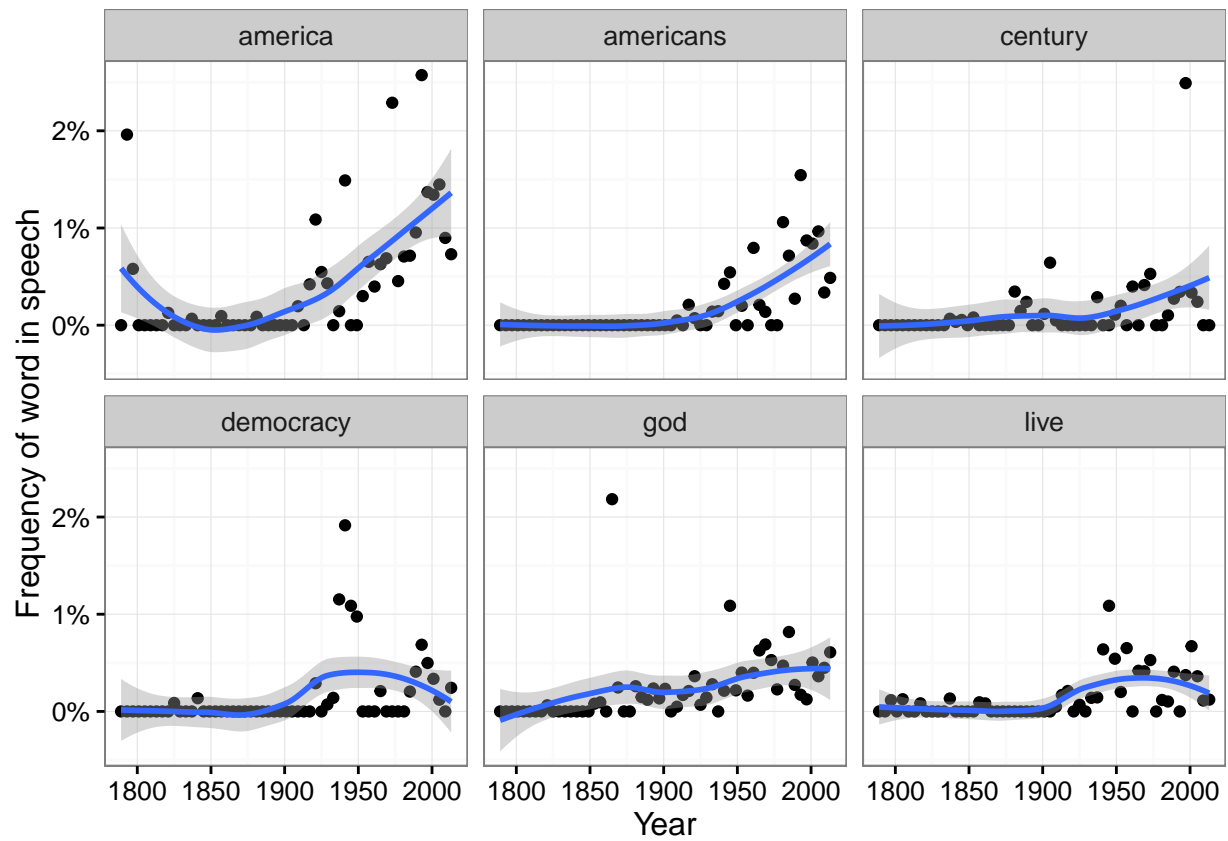
```



We can also use the ggplot2 package to display the top 6 terms that have changed in frequency over time.

```
library(scales)

models %>%
  top_n(6, abs(estimate)) %>%
  inner_join(inaug_freq) %>%
  ggplot(aes(Year, percent)) +
  geom_point() +
  geom_smooth() +
  facet_wrap(~ word) +
  scale_y_continuous(labels = percent_format()) +
  ylab("Frequency of word in speech")
```



Chapter 6

Topic Modeling

Topic modeling is a method for unsupervised classification of documents, by modeling each document as a mixture of topics and each topic as a mixture of words. Latent Dirichlet allocation is a particularly popular method for fitting a topic model.

We can use tidy text principles, as described in the main vignette, to approach topic modeling using consistent and effective tools. In particular, we'll be using tidying functions for LDA objects from the `topicmodels` package.

6.1 Can we tell the difference between Dickens, Wells, Verne, and Austen?

Suppose a vandal has broken into your study and torn apart four of your books:

- *Great Expectations* by Charles Dickens
- *The War of the Worlds* by H.G. Wells
- *Twenty Thousand Leagues Under the Sea* by Jules Verne
- *Pride and Prejudice* by Jane Austen

This vandal has torn the books into individual chapters, and left them in one large pile. How can we restore these disorganized chapters to their original books?

6.2 Setup

We'll retrieve four books using the `gutenbergr` package:

```
library(dplyr)
library(gutenbergr)

titles <- c("Twenty Thousand Leagues under the Sea", "The War of the Worlds",
           "Pride and Prejudice", "Great Expectations")

books <- gutenberg_works(title %in% titles) %>%
  gutenberg_download(meta_fields = "title")

books
```

```
## # A tibble: 51,663 x 3
```

```
##      gutenber_id                text                title
##      <int>                  <chr>                  <chr>
## 1         36      The War of the Worlds The War of the Worlds
## 2         36                        The War of the Worlds
## 3         36      by H. G. Wells [1898] The War of the Worlds
## 4         36                        The War of the Worlds
## 5         36                        The War of the Worlds
## 6         36      But who shall dwell in these worlds if they be The War of the Worlds
## 7         36      inhabited? . . . Are we or they Lords of the The War of the Worlds
## 8         36      World? . . . And how are all things made for man?-- The War of the Worlds
## 9         36      KEPLER (quoted in The Anatomy of Melancholy) The War of the Worlds
## 10        36                        The War of the Worlds
## # ... with 51,653 more rows
```

As pre-processing, we divide these into chapters, use `tidytext`'s `unnest_tokens` to separate them into words, then remove `stop_words`. We're treating every chapter as a separate "document", each with a name like `Great Expectations_1` or `Pride and Prejudice_11`.

```
library(tidytext)
library(stringr)
library(tidyr)

by_chapter <- books %>%
  group_by(title) %>%
  mutate(chapter = cumsum(str_detect(text, regex("^chapter ", ignore_case = TRUE)))) %>%
  ungroup() %>%
  filter(chapter > 0)

by_chapter_word <- by_chapter %>%
  unite(title_chapter, title, chapter) %>%
  unnest_tokens(word, text)

word_counts <- by_chapter_word %>%
  anti_join(stop_words) %>%
  count(title_chapter, word, sort = TRUE) %>%
  ungroup()

word_counts
```

```
## # A tibble: 104,721 x 3
##       title_chapter  word      n
##       <chr>      <chr> <int>
## 1 Great Expectations_1  sir     13
## 2 Great Expectations_1 church     7
## 3 Great Expectations_1 looked     7
## 4 Great Expectations_1  pip     7
## 5 Great Expectations_1  river     6
## 6 Great Expectations_1 tilted     6
## 7 Great Expectations_1  black     5
## 8 Great Expectations_1  head     5
## 9 Great Expectations_1  live     5
## 10 Great Expectations_1 mother     5
## # ... with 104,711 more rows
```

6.3 Latent Dirichlet Allocation with the topicmodels package

Right now this data frame is in a tidy form, with one-term-per-document-per-row. However, the `topicmodels` package requires a `DocumentTermMatrix` (from the `tm` package). As described in this vignette, we can cast a one-token-per-row table into a `DocumentTermMatrix` with `tidytext`'s `cast_dtm`:

```
chapters_dtm <- word_counts %>%
  cast_dtm(title_chapter, word, n)
```

```
chapters_dtm
```

```
## <<DocumentTermMatrix (documents: 193, terms: 18215)>>
## Non-/sparse entries: 104721/3410774
## Sparsity          : 97%
## Maximal term length: 19
## Weighting          : term frequency (tf)
```

Now we are ready to use the `topicmodels` package to create a four topic LDA model.

```
library(topicmodels)
chapters_lda <- LDA(chapters_dtm, k = 4, control = list(seed = 1234))
chapters_lda
```

```
## A LDA_VEM topic model with 4 topics.
```

(In this case we know there are four topics because there are four books; in practice we may need to try a few different values of k).

Now `tidytext` gives us the option of *returning* to a tidy analysis, using the `tidy` and `augment` verbs borrowed from the `broom` package. In particular, we start with the `tidy` verb.

```
chapters_lda_td <- tidy(chapters_lda)
chapters_lda_td
```

```
## # A tibble: 72,860 x 3
##   topic  term      beta
##   <int> <chr>    <dbl>
## 1     1   sir 2.084216e-03
## 2     2   sir 4.985743e-03
## 3     3   sir 2.483177e-03
## 4     4   sir 2.935244e-04
## 5     1 church 2.936533e-04
## 6     2 church 7.641174e-05
## 7     3 church 7.494668e-04
## 8     4 church 7.008334e-04
## 9     1 looked 1.950223e-03
## 10    2 looked 1.533431e-03
## # ... with 72,850 more rows
```

Notice that this has turned the model into a one-topic-per-term-per-row format. For each combination the model has β , the probability of that term being generated from that topic.

We could use `dplyr`'s `top_n` to find the top 5 terms within each topic:

```
top_terms <- chapters_lda_td %>%
  group_by(topic) %>%
  top_n(5, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)
```

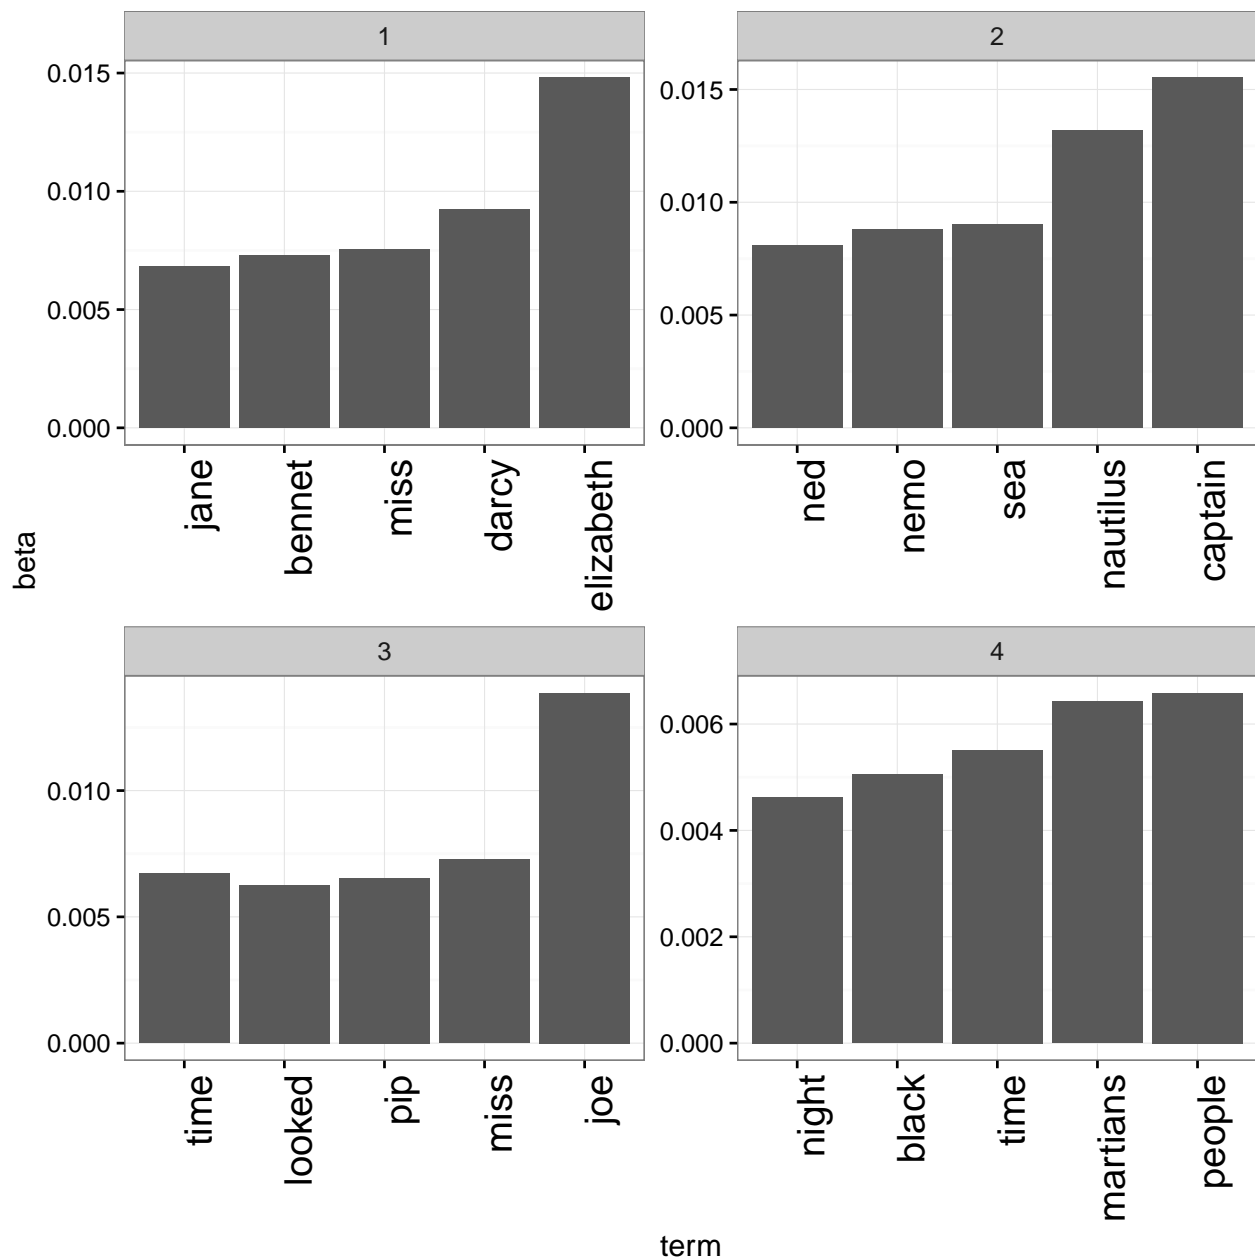
```
top_terms
```

```
## # A tibble: 20 x 3
##   topic      term      beta
##   <int>    <chr>    <dbl>
## 1      1 elizabeth 0.014806162
## 2      1 darcy    0.009250751
## 3      1 miss     0.007560066
## 4      1 bennet   0.007291477
## 5      1 jane     0.006845060
## 6      2 captain  0.015526574
## 7      2 nautilus 0.013176023
## 8      2 sea      0.009007086
## 9      2 nemo     0.008792461
## 10     2 ned      0.008108322
## 11     3 joe      0.013866397
## 12     3 miss     0.007262017
## 13     3 time     0.006709286
## 14     3 pip      0.006541889
## 15     3 looked  0.006262112
## 16     4 people   0.006584429
## 17     4 martians 0.006419187
## 18     4 time     0.005508969
## 19     4 black    0.005051646
## 20     4 night    0.004629963
```

This model lends itself to a visualization:

```
library(ggplot2)
theme_set(theme_bw())

top_terms %>%
  mutate(term = reorder(term, beta)) %>%
  ggplot(aes(term, beta)) +
  geom_bar(stat = "identity") +
  facet_wrap(~ topic, scales = "free") +
  theme(axis.text.x = element_text(size = 15, angle = 90, hjust = 1))
```

These topics are pretty clearly associated with the four books! There’s no question that the topic of “nemo”, “sea”, and “nautilus” belongs to *Twenty Thousand Leagues Under the Sea*, and that “jane”, “darcy”, and “elizabeth” belongs to *Pride and Prejudice*. We see “pip” and “joe” from *Great Expectations* and “martians”, “black”, and “night” from *The War of the Worlds*.

6.4 Per-document classification

Each chapter was a “document” in this analysis. Thus, we may want to know which topics are associated with each document. Can we put the chapters back together in the correct books?

```
chapters_lda_gamma <- tidy(chapters_lda, matrix = "gamma")
chapters_lda_gamma
```

```
## # A tibble: 772 x 3
```

```
##           document topic      gamma
##           <chr> <int>      <dbl>
## 1 Great Expectations_1      1 3.012893e-05
## 2 Great Expectations_10     1 2.145733e-05
## 3 Great Expectations_11     1 1.064330e-05
## 4 Great Expectations_12     1 2.477705e-05
## 5 Great Expectations_13     1 1.954208e-05
## 6 Great Expectations_14     1 7.234765e-05
## 7 Great Expectations_15     1 1.327335e-05
## 8 Great Expectations_16     1 2.951203e-05
## 9 Great Expectations_17     1 1.954208e-05
## 10 Great Expectations_18    1 1.174153e-05
## # ... with 762 more rows
```

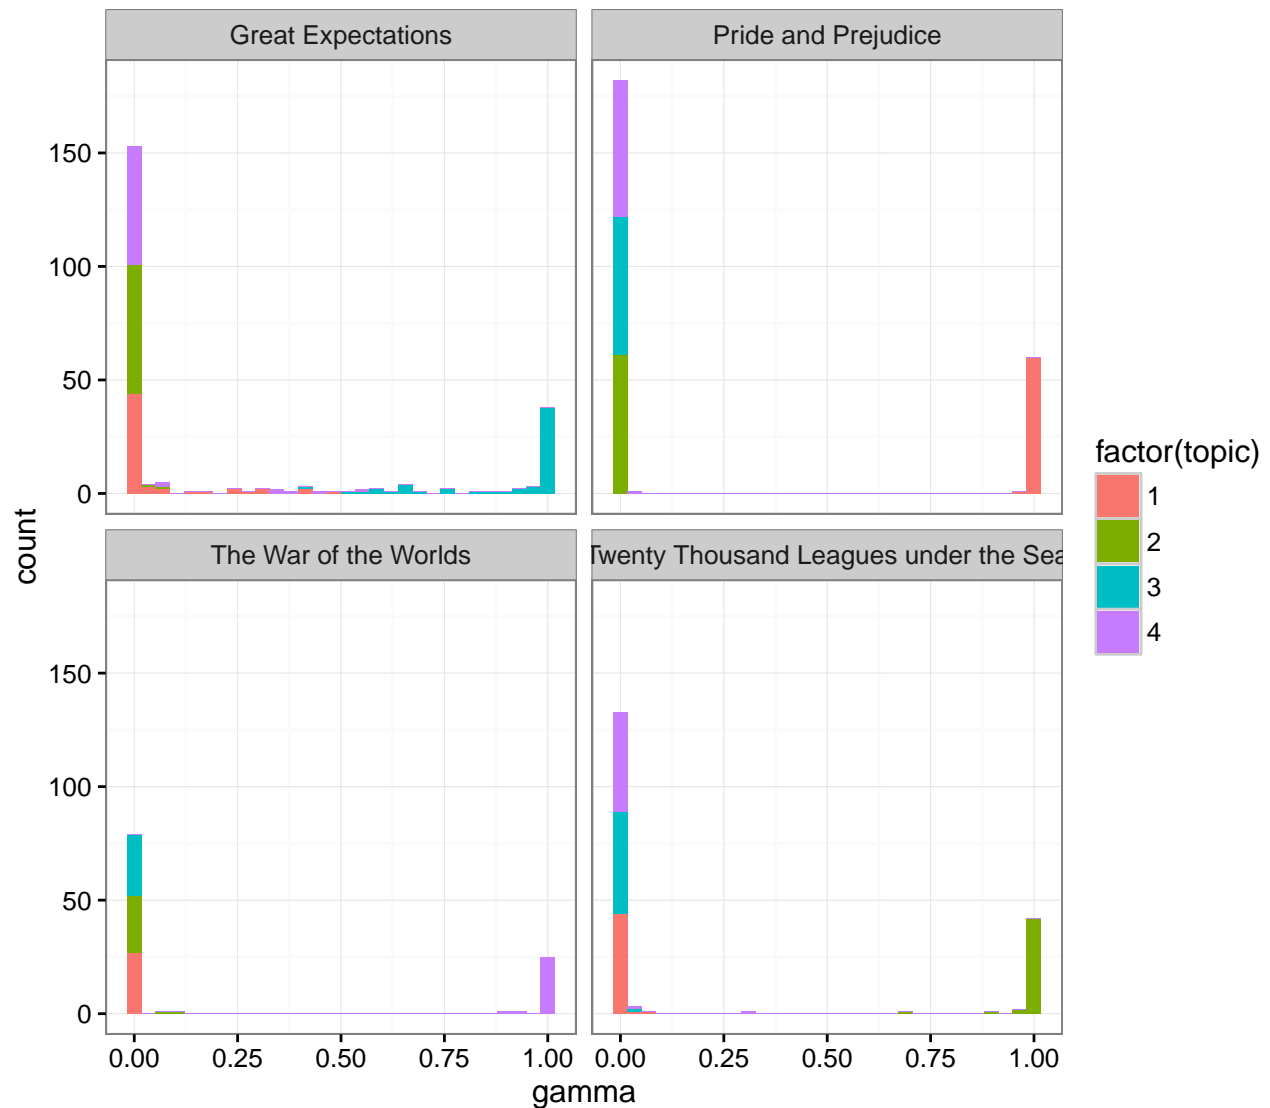
Setting `matrix = "gamma"` returns a tidied version with one-document-per-topic-per-row. Now that we have these document classifications, we can see how well our unsupervised learning did at distinguishing the four books. First we re-separate the document name into title and chapter:

```
chapters_lda_gamma <- chapters_lda_gamma %>%
  separate(document, c("title", "chapter"), sep = "_", convert = TRUE)
chapters_lda_gamma
```

```
## # A tibble: 772 x 4
##           title chapter topic      gamma
## *      <chr>    <int> <int>      <dbl>
## 1 Great Expectations      1      1 3.012893e-05
## 2 Great Expectations     10      1 2.145733e-05
## 3 Great Expectations     11      1 1.064330e-05
## 4 Great Expectations     12      1 2.477705e-05
## 5 Great Expectations     13      1 1.954208e-05
## 6 Great Expectations     14      1 7.234765e-05
## 7 Great Expectations     15      1 1.327335e-05
## 8 Great Expectations     16      1 2.951203e-05
## 9 Great Expectations     17      1 1.954208e-05
## 10 Great Expectations     18      1 1.174153e-05
## # ... with 762 more rows
```

Then we examine what fraction of chapters we got right for each:

```
ggplot(chapters_lda_gamma, aes(gamma, fill = factor(topic))) +
  geom_histogram() +
  facet_wrap(~ title, nrow = 2)
```



We notice that almost all of the chapters from *Pride and Prejudice*, *War of the Worlds*, and *Twenty Thousand Leagues Under the Sea* were uniquely identified as a single topic each.

```
chapter_classifications <- chapters_lda_gamma %>%
  group_by(title, chapter) %>%
  top_n(1, gamma) %>%
  ungroup() %>%
  arrange(gamma)

chapter_classifications
```

```
## # A tibble: 193 x 4
##       title chapter topic  gamma
##   <chr>   <int> <int>  <dbl>
## 1 Great Expectations    23     3 0.5187562
## 2 Great Expectations    54     4 0.5522191
## 3 Great Expectations    56     3 0.5523846
## 4 Great Expectations    37     3 0.5724660
## 5 Great Expectations    55     3 0.5813702
```

```
## 6 Great Expectations      46      3 0.6067932
## 7 Great Expectations      25      3 0.6405668
## 8 Great Expectations      21      3 0.6443639
## 9 Great Expectations      20      3 0.6457449
## 10 Great Expectations     53      3 0.6636836
## # ... with 183 more rows
```

We can determine this by finding the consensus book for each, which we note is correct based on our earlier visualization:

```
book_topics <- chapter_classifications %>%
  count(title, topic) %>%
  top_n(1, n) %>%
  ungroup() %>%
  transmute(consensus = title, topic)
```

```
book_topics
```

```
## # A tibble: 4 x 2
##               consensus topic
##               <chr> <int>
## 1 Great Expectations      3
## 2 Pride and Prejudice      1
## 3 The War of the Worlds     4
## 4 Twenty Thousand Leagues under the Sea 2
```

Then we see which chapters were misidentified:

```
chapter_classifications %>%
  inner_join(book_topics, by = "topic") %>%
  count(title, consensus)
```

```
## Source: local data frame [5 x 3]
## Groups: title [?]
##
##               title                consensus      n
##               (chr)                (chr) (int)
## 1 Great Expectations Great Expectations  58
## 2 Great Expectations The War of the Worlds   1
## 3 Pride and Prejudice Pride and Prejudice  61
## 4 The War of the Worlds The War of the Worlds  27
## 5 Twenty Thousand Leagues under the Sea Twenty Thousand Leagues under the Sea 46
```

We see that only a few chapters from *Great Expectations* were misclassified. Not bad for unsupervised clustering!

6.4.1 By word assignments: `augment`

One important step in the topic modeling expectation-maximization algorithm is assigning each word in each document to a topic. The more words in a document are assigned to that topic, generally, the more weight (γ) will go on that document-topic classification.

We may want to take the original document-word pairs and find which words in each document were assigned to which topic. This is the job of the `augment` verb.

```
assignments <- augment(chapters_lda, data = chapters_dtm)
```

We can combine this with the consensus book titles to find which words were incorrectly classified.

```
assignments <- assignments %>%
  separate(document, c("title", "chapter"), sep = "_", convert = TRUE) %>%
  inner_join(book_topics, by = c(".topic" = "topic"))
```

```
assignments
```

```
## # A tibble: 104,721 x 6
##           title chapter term count .topic consensus
##           <chr>   <int> <chr> <dbl> <dbl>      <chr>
## 1 Great Expectations      1  sir    13      3 Great Expectations
## 2 Great Expectations     10  sir     1      3 Great Expectations
## 3 Great Expectations     11  sir     2      3 Great Expectations
## 4 Great Expectations     18  sir     3      3 Great Expectations
## 5 Great Expectations     19  sir    10      3 Great Expectations
## 6 Great Expectations     20  sir     7      3 Great Expectations
## 7 Great Expectations     23  sir     1      3 Great Expectations
## 8 Great Expectations     25  sir     6      3 Great Expectations
## 9 Great Expectations     26  sir     2      3 Great Expectations
## 10 Great Expectations     27  sir    18      3 Great Expectations
## # ... with 104,711 more rows
```

We can, for example, create a “confusion matrix” using `dplyr`’s `count` and `tidyr`’s `spread`:

```
assignments %>%
  count(title, consensus, wt = count) %>%
  spread(consensus, n, fill = 0)
```

```
## Source: local data frame [4 x 5]
## Groups: title [4]
##
##           title Great Expectations Pride and Prejudice
##           (chr)              (dbl)              (dbl)
## 1 Great Expectations          51043              2146
## 2 Pride and Prejudice           0              37241
## 3 The War of the Worlds         0               0
## 4 Twenty Thousand Leagues under the Sea 3              31
## The War of the Worlds Twenty Thousand Leagues under the Sea
##           (dbl)              (dbl)
## 1 2341              38
## 2 1                 0
## 3 22492             76
## 4 303              39297
```

We notice that almost all the words for *Pride and Prejudice*, *Twenty Thousand Leagues Under the Sea*, and *War of the Worlds* were correctly assigned, while *Great Expectations* had a fair amount of misassignment.

What were the most commonly mistaken words?

```
wrong_words <- assignments %>%
  filter(title != consensus)
```

```
wrong_words
```

```
## # A tibble: 3,701 x 6
##           title chapter term count .topic consensus
##           <chr>   <int> <chr> <dbl> <dbl>      <chr>
## 1 Great Expectations     46  river     5      4 The War of the Worlds
```

```
## 2          Great Expectations    54 river    12    4 The War of the Worlds
## 3          Great Expectations    20 black     3    4 The War of the Worlds
## 4          Great Expectations    46 black     2    4 The War of the Worlds
## 5          Great Expectations    53 black     1    4 The War of the Worlds
## 6          Great Expectations    54 black     2    4 The War of the Worlds
## 7 Twenty Thousand Leagues under the Sea    18 black     1    4 The War of the Worlds
## 8          Great Expectations    23 live      1    1 Pride and Prejudice
## 9          Great Expectations    54 live      1    4 The War of the Worlds
## 10         Great Expectations    25 mother    2    1 Pride and Prejudice
## # ... with 3,691 more rows
```

```
wrong_words %>%
  count(title, consensus, term, wt = count) %>%
  ungroup() %>%
  arrange(desc(n))
```

```
## # A tibble: 3,027 x 4
##       title          consensus      term      n
##       <chr>          <chr>      <chr> <dbl>
## 1 Great Expectations The War of the Worlds    boat    30
## 2 Great Expectations Pride and Prejudice skiffins  25
## 3 Great Expectations The War of the Worlds    tide    25
## 4 Great Expectations The War of the Worlds    lay     19
## 5 Great Expectations The War of the Worlds    jack    18
## 6 Great Expectations The War of the Worlds    water    18
## 7 Great Expectations Pride and Prejudice    lady    17
## 8 Great Expectations The War of the Worlds    river    17
## 9 Great Expectations The War of the Worlds    barley   16
## 10 Great Expectations The War of the Worlds    galley   16
## # ... with 3,017 more rows
```

Notice the word “flopson” here; these wrong words do not necessarily appear in the novels they were misassigned to. Indeed, we can confirm “flopson” appears only in *Great Expectations*:

```
word_counts %>%
  filter(word == "flopson")
```

```
## # A tibble: 3 x 3
##       title_chapter  word      n
##       <chr>      <chr> <int>
## 1 Great Expectations_22 flopson    10
## 2 Great Expectations_23 flopson     7
## 3 Great Expectations_33 flopson     1
```

The algorithm is stochastic and iterative, and it can accidentally land on a topic that spans multiple books.

Chapter 7

Tidying word2vec Models from the glove Package

TODO: still a lot of work to be done on the methods as well as the chapter, may or may not make it in

Chapter 8

Predicting ratings from text in the Yelp food reviews dataset

Intro goes here

8.1 Setup

I've downloaded the `yelp_dataset_challenge_academic_dataset` folder from [here](#).^[^termsofuse] First I read and process them.

```
library(readr)
library(dplyr)

# You may have used the built-in readLines before, but read_lines from
# readr is faster for large files

# we're reading only 100,000 in this example
# you can try it with the full dataset too, it's just a little slower!
# in the final version of the book we're probably going to read all, it
# just makes this chapter take a while to compile

infile <- "~/Downloads/yelp_dataset_challenge_academic_dataset/yelp_academic_dataset_review.json"
review_lines <- read_lines(infile, n_max = 100000)

library(stringr)

# Each line is a JSON object- the fastest way to process is to combine into a
# single JSON string and use jsonlite::fromJSON
reviews_combined <- str_c("[", str_c(review_lines, collapse = ", "), "]")

reviews <- jsonlite::fromJSON(reviews_combined) %>%
  jsonlite::flatten() %>%
  tbl_df()

reviews

## # A tibble: 100,000 x 10
##           user_id           review_id stars      date
##           <chr>           <chr> <int>   <chr>
```

```
## 1  PUFPaY9KxDacGqfsorJp3Q Ya85v4eqdd6k90d8HbQjyA      4 2012-08-01
## 2  Iu6AxdBYGR4A0wspR9BYHA KpVLNJ21_4wbYNctrOwWdQ      5 2014-02-13
## 3  auESFwWvW42h6alXgFxAxQ fFS0GV46Yxuwbr3fHNUZig      5 2015-10-31
## 4  uK8tzra0p4M5u3uYrqIBXg Di3exaUCFNw1V4kSNW5pgA      5 2013-11-08
## 5  I_47G-R2_egp7ME5u_ltew 0Lua2-PbqEQMjD9r89-asw      3 2014-03-29
## 6  PP_xoMSYlGr2pb67BbqBdA 7N9j5YbBHBW6qguE5DAeyA      1 2014-10-29
## 7  JPPhyFE-UE453zA6K0TVgw mJCJR33jvUNT41iJCxDU_g      4 2014-11-28
## 8  2d5HeDvZTDUNVog_WuUpSg Ieh3kfZ-5J9pLju4JiQDvQ      5 2014-02-27
## 9  BShxMIUwaJS378xcrcz4NmG PU280oBSHpZLkYGCmNxlmg      5 2015-06-16
## 10 fhNx0MwWTipzj08A9LFe8Q XsA6AojkWjOHA4FmuAb8XQ      3 2012-08-19
## # ... with 99,990 more rows, and 6 more variables: text <chr>, type <chr>, business_id <chr>,
## #   votes.funny <int>, votes.useful <int>, votes.cool <int>
```

8.2 Tidy sentiment analysis

Right now, there is one row for each review. To analyze in the tidy text framework, we need to use the `unnest_tokens` function and turn this into one-row-per-term-per-document:

```
library(tidytext)

review_words <- reviews %>%
  select(review_id, business_id, stars, text) %>%
  unnest_tokens(word, text) %>%
  filter(!word %in% stop_words$word,
         str_detect(word, "[a-z]+$"))

review_words

## # A tibble: 3,971,444 x 4
##       review_id      business_id stars      word
##       <chr>          <chr> <int>    <chr>
## 1 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4    hoagie
## 2 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4 institution
## 3 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4    walking
## 4 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4   throwback
## 5 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4        ago
## 6 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4   fashioned
## 7 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4        menu
## 8 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4        board
## 9 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4       booths
## 10 Ya85v4eqdd6k90d8HbQjyA 5UmKMjUEUNdYWqANhGckJw      4   selection
## # ... with 3,971,434 more rows
```

Notice that there is now one-row-per-term-per-document: the In this cleaning process we’ve also removed “stopwords” (such as “I”, “the”, “and”, etc), and removing things things that are formatting (e.g. “—”) rather than a word.

Now I’m going to do sentiment analysis on each review. We’ll use the AFINN lexicon, which provides a positivity score for each word, from -5 (most negative) to 5 (most positive).

```
AFINN <- sentiments %>%
  filter(lexicon == "AFINN") %>%
  select(word, afinn_score = score)

AFINN
```

```
## # A tibble: 2,476 x 2
##       word afinn_score
##       <chr>      <int>
## 1  abandon      -2
## 2  abandoned    -2
## 3  abandons     -2
## 4  abducted     -2
## 5  abduction    -2
## 6  abductions   -2
## 7   abhor       -3
## 8  abhorred     -3
## 9  abhorrent    -3
## 10 abhors       -3
## # ... with 2,466 more rows
```

Now as described in this post, our sentiment analysis is just an inner-join operation followed by a summary:

```
reviews_sentiment <- review_words %>%
  inner_join(AFINN, by = "word") %>%
  group_by(review_id, stars) %>%
  summarize(sentiment = mean(afinn_score))


reviews_sentiment
```

```
## Source: local data frame [93,947 x 3]
## Groups: review_id [?]
##
##       review_id stars  sentiment
##       (chr) (int)    (dbl)
## 1  __-r0eC3hZlaejvuliC8zQ      5  4.0000000
## 2  __56FUEaW57kZEm560Zk7w      5  0.8333333
## 3  __6t0xx2VcvGR02d2ILkuw      5  1.7500000
## 4  __77nP3Nf1wsGz5HPs2hdw      5  1.6000000
## 5  __B5KInsYxFKIHKXAS6_rA      1 -2.0000000
## 6  __BIQ3tcFZg6_PpdadEfLQ      4  1.6000000
## 7  __DK9Vsmyo0zJQhI15cbg      1 -2.1000000
## 8  __ELCJ0wzDM2QNRfVUq26Q      5  3.5000000
## 9  __esH_kgJZeS8k3i6HaG7Q      5  0.2142857
## 10 __GXnNfKFLqFhMtpCTTT2g      3  0.8750000
## ..      ...      ...      ...
```

Now we can see how our estimates did!

```
library(ggplot2)
theme_set(theme_bw())

ggplot(reviews_sentiment, aes(stars, sentiment, group = stars)) +
  geom_boxplot() +
  ylab("Average sentiment score")
```



Well, it's a good start! Our sentiment scores are correlated with positivity ratings. But we do see that there's a large amount of prediction error- some 5-star reviews have a highly negative sentiment score, and vice versa.

8.3 Which words are positive or negative?

We're interested in analyzing the properties of words. Which are suggestive of positive reviews, and which are negative? To do this, we'll create a per-word summary.

```
review_words_counted <- review_words %>%
  count(review_id, business_id, stars, word) %>%
  ungroup()
```

```
review_words_counted
```

```
## # A tibble: 3,405,173 x 5
##   review_id business_id stars word n
##   <chr>      <chr> <int> <chr> <int>
## 1 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 batter 1
## 2 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 chips 3
## 3 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 compares 1
## 4 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 fashioned 1
## 5 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 filleted 1
## 6 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 fish 4
## 7 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 fries 1
## 8 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 frozen 1
## 9 ___XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w 5 greenlake 1
```

```
## 10 __XYEos-RIkPsQwplRYyw YxMnfznT3eYya0YV37tE8w      5      hand      1
## # ... with 3,405,163 more rows
```

```
word_summaries <- review_words_counted %>%
  group_by(word) %>%
  summarize(reviews = n(),
            uses = sum(n),
            average_stars = mean(stars)) %>%
  ungroup()
```

```
word_summaries
```

```
## # A tibble: 73,816 x 4
##       word reviews  uses average_stars
##       <chr>    <int> <int>         <dbl>
## 1  a'boiling      1      1           4.00
## 2  a'fare         1      1           4.00
## 3  a'ight         2      2           1.50
## 4  a'la           2      2           4.50
## 5  a'll           1      1           1.00
## 6  a'lyce         1      2           5.00
## 7  a'more         2      2           5.00
## 8  a'orange       1      1           5.00
## 9  a'prowling     1      1           3.00
## 10 aa           20     23           3.25
## # ... with 73,806 more rows
```

We can start by looking only at words that appear in at least 100 (out of 100000) reviews. This makes sense both because words that appear more rarely will have a noisier measurement (a few good or bad reviews could shift the balance), and because they're less likely to be useful in classifying future reviews or text.

```
word_summaries_filtered <- word_summaries %>%
  filter(reviews >= 100)
```

```
word_summaries_filtered
```

```
## # A tibble: 4,465 x 4
##       word reviews  uses average_stars
##       <chr>    <int> <int>         <dbl>
## 1    aaa       100   145      3.780000
## 2  ability     210   215      3.580952
## 3 absolute     589   600      3.755518
## 4 absolutely  3195  3401      3.812520
## 5    ac        306   420      3.058824
## 6  accent     112   115      3.446429
## 7  accept     350   370      3.060000
## 8 acceptable  313   319      2.645367
## 9  accepted   162   167      3.030864
## 10 access     530   588      3.541509
## # ... with 4,455 more rows
```

What were the most positive and negative words?

```
word_summaries_filtered %>%
  arrange(desc(average_stars))
```

```
## # A tibble: 4,465 x 4
##       word reviews  uses average_stars
```


```
##           <chr>    <int> <int>         <dbl>
## 1      exceeded      160   161      4.675000
## 2    knowledgeable      371   374      4.630728
## 3   compassionate      112   115      4.625000
## 4       exquisite      108   112      4.601852
## 5        chihuly      117   151      4.547009
## 6        treasure      152   159      4.546053
## 7    compliments      212   215      4.542453
## 8           gem       982   997      4.518330
## 9    botanical       184   241      4.510870
## 10  trustworthy       105   105      4.495238
## # ... with 4,455 more rows
```

```
word_summaries_filtered %>%
  arrange(average_stars)
```

```
## # A tibble: 4,465 x 4
##           word reviews  uses average_stars
##           <chr>    <int> <int>         <dbl>
## 1    incompetent      156   167      1.423077
## 2 unprofessional      362   383      1.447514
## 3      disgusted      124   126      1.475806
## 4        rudely      167   179      1.520958
## 5          lied      154   177      1.532468
## 6        refund      390   497      1.558974
## 7        refused      455   507      1.610989
## 8   unacceptable      201   203      1.616915
## 9          worst     2433  2653      1.619400
## 10        denied       100   111      1.650000
## # ... with 4,455 more rows
```

Makes a lot of sense! We can also plot positivity by frequency:

```
ggplot(word_summaries_filtered, aes(reviews, average_stars)) +
  geom_point() +
  geom_text(aes(label = word), check_overlap = TRUE, vjust = 1, hjust = 1) +
  scale_x_log10() +
  geom_hline(yintercept = mean(reviews$stars), color = "red", lty = 2) +
  xlab("# of reviews") +
  ylab("Average Stars")
```



Note that some of the most common words (e.g. “food”) are pretty neutral. There are some common words that are pretty positive (e.g. “amazing”, “awesome”) and others that are pretty negative (“bad”, “told”).

8.4 Comparing to sentiment analysis

When we perform sentiment analysis, we're often comparing to a pre-existing lexicon, one that was developed.

The tidytext package also comes with several tidy sentiment analysis lexicons:

sentiments

```
## # A tibble: 23,165 x 4
##       word sentiment  lexicon score
##       <chr>      <chr>    <chr> <int>
## 1    abacus      trust     nrc    NA
## 2   abandon     fear     nrc    NA
## 3   abandon negative     nrc    NA
## 4   abandon sadness     nrc    NA
## 5 abandoned    anger     nrc    NA
## 6 abandoned    fear     nrc    NA
## 7 abandoned negative     nrc    NA
## 8 abandoned    sadness  nrc    NA
## 9 abandonment  anger     nrc    NA
## 10 abandonment fear     nrc    NA
## # ... with 23,155 more rows
```

We might expect that more positive words are associated with higher star reviews. Does this hold? We can combine and compare the two datasets with `inner_join`.


```
words_afinn <- word_summaries_filtered %>%
  inner_join(AFINN)
```

```
words_afinn
```

```
## # A tibble: 520 x 5
```


```
##       word reviews uses average_stars afinn_score
##       <chr>   <int> <int>      <dbl>      <int>
## 1    ability    210   215      3.580952         2
## 2     accept    350   370      3.060000         1
## 3   accepted    162   167      3.030864         1
## 4   accident    213   239      3.629108        -2
## 5 accidentally  152   152      3.348684        -2
## 6     active    109   115      3.981651         1
## 7   adequate    290   304      3.262069         1
## 8     admit    740   754      3.666216        -1
## 9   admitted    111   118      2.225225        -1
## 10  adorable    255   266      4.250980         3
## # ... with 510 more rows
```

```
ggplot(words_afinn, aes(afinn_score, average_stars, group = afinn_score)) +
  geom_boxplot() +
  xlab("AFINN score of word") +
  ylab("Average stars of reviews with this word")
```



Just like in our per-review predictions, there's a very clear trend. AFINN sentiment analysis works, at least a little bit!


But we may want to see some of those details. Which positive/negative words were most successful in predicting a positive/negative review, and which broke the trend?



```
## mapping: x = x
## geom_blank: na.rm = FALSE
## stat_identity: na.rm = FALSE
## position_identity
```

For example, we can see that most curse words have an AFINN score of -4, and that while some words, like “wtf”, successfully predict a negative review, others, like “damn”, are often positive. (They’re likely part of “damn good”, or something similar). Some of the words that AFINN most underestimated included “die” (“the pork chops are to **die** for!”), and one of the words it most overestimated was “joke” (“the service is a complete **joke**!”).

One other way we could look at mis



Chapter 9

Some analysis goes here

I don't know what will go here, but I'd like to have one more analysis that touches on all of the areas of tidy text analysis. If we can't find one we'll skip it!